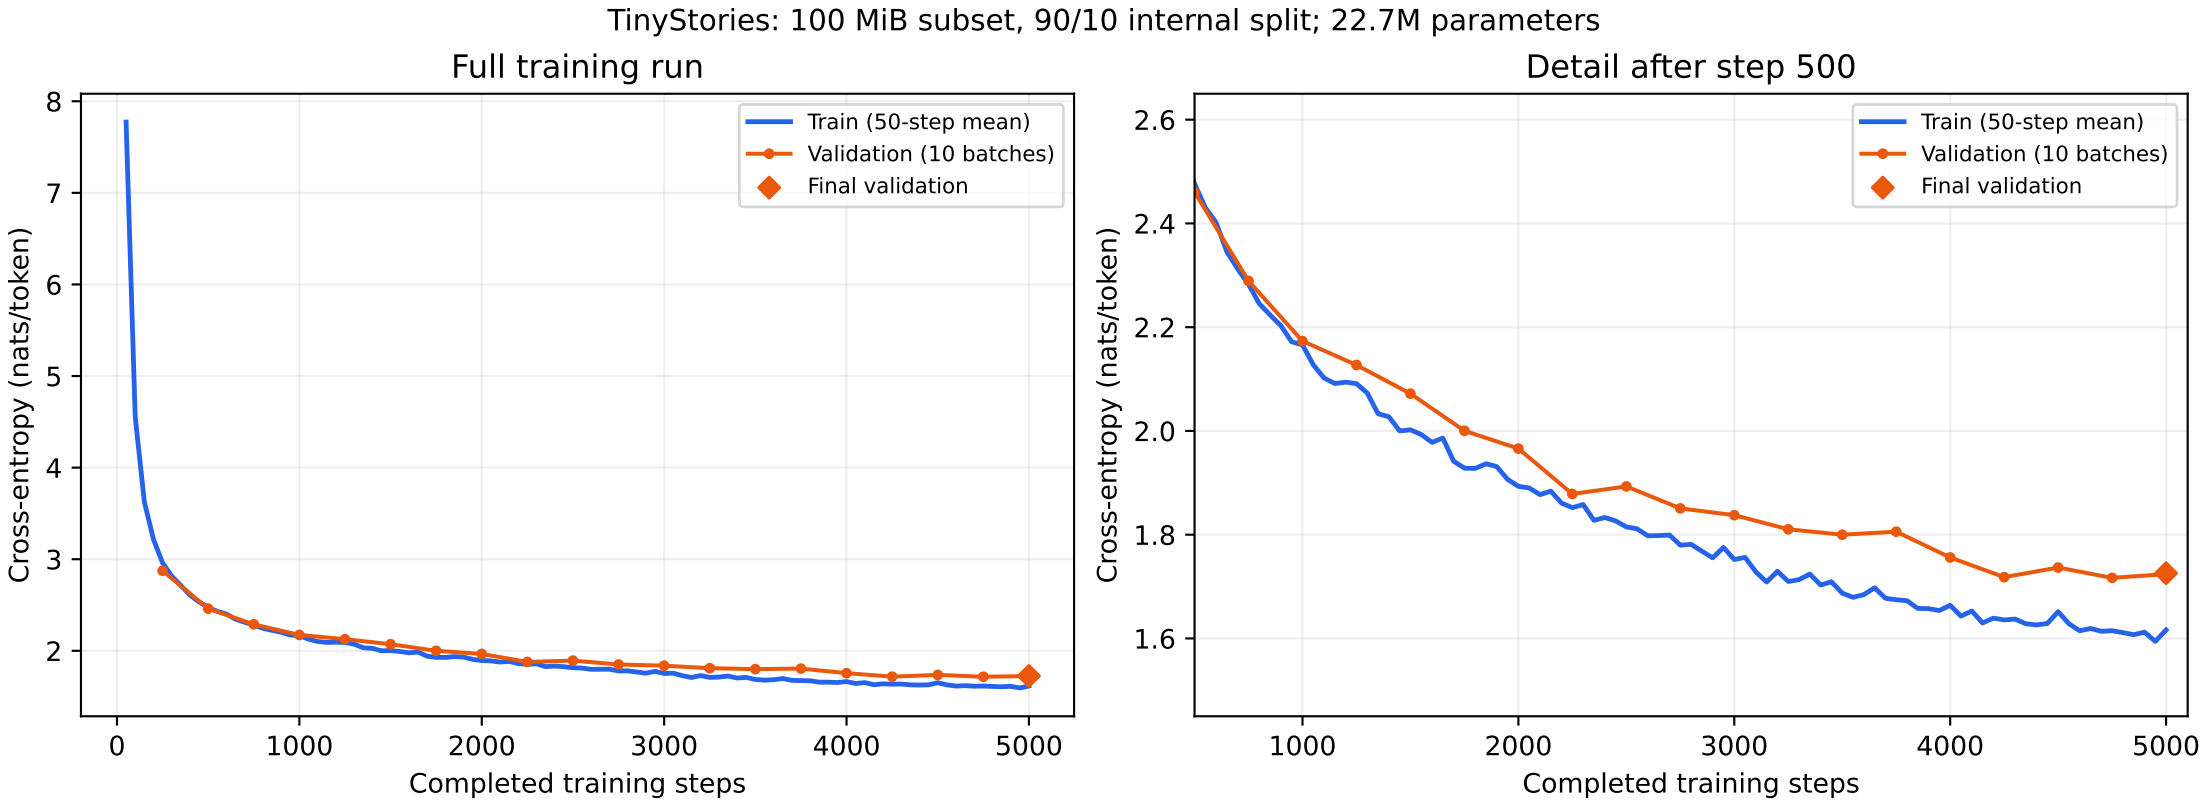
TinyStories: 100 MiB subset, 90/10 internal split; 22.7M parameters
Full training run
Detail after step 500
Train (50-step mean)
Validation (10 batches)
Final validation
Cross-entropy (nats/token)
Completed training steps
0
1000
2000
3000
4000
5000
2
3
4
5
6
7
8
1.6
1.8
2.0
2.2
2.4
2.6